OutOfRole: Build Plan

The Debugging Table · IBM Z Datathon 2026 · October 17 to 18

What we're building

OutOfRole spots insiders while they are still active, explains each alert with checked evidence, and trims everyone's access to what they actually need. It all runs on one LinuxONE server, and people appear only as pseudonyms (U-017).

1. **Early warning (the headline).** Replay the logs day by day. Compare each person with their own normal weeks; the 20 most unusual people form a daily review list.
2. **Granite case files.** For each person on the list, code picks the evidence ("new to this person, rarest first"). Granite writes the explanation citing it, and code checks every number before anyone sees it.
3. **Least-privilege advisor (the recommender).** A recommender predicts which websites each person will legitimately need next. Those are pre-approved; anything else needs a request. Sites nobody in the company has ever visited are blocked.

Pitch line: "Most insider tools find the culprit after they've left. Ours put 80 of 99 insiders on a 20-name daily list while they were still active, before about half of their malicious actions. Granite, on the same server, writes a case file citing the real evidence. A recommender trims everyone's access to what they actually need."

Numbers we tested

CERT data. Settings were frozen on r4.2, then r5.2 was run unchanged. Each number was reproduced with independent code.

Claim	r4.2 (70 insiders)	r5.2 (99 insiders)
Insiders on the daily list before their last malicious action	**59**	**80**
Same, with a simple 28-day personal baseline	33	28
Same, by chance	about 3	about 3
Share of malicious actions that happen after the flag	54%	48%
New names per week for the analyst	about 5	about 5
Flagged insiders with a real malicious event in their top 3 evidence items	98%	98%
Next new websites caught by the recommender's top 20 (popularity in brackets)	21% (14.5%)	19% (13%)
Access requests per person per month: history only, then with the recommender	2.3 to 1.8	3.1 to 2.5
Wikileaks and keylogger visits blocked by the new-to-company rule	100%	100%

By attack type:

- Leakers are flagged the same day.
- Job-hunters and snoopers are flagged after 1.5 to 3 weeks.
- Two-day sysadmin sabotage is mostly missed: 4 of 10 on r4.2, 2 of 10 on r5.2.

Step by step: what each test changed

"Found" = insiders expected in the top 70. Full detail: the notebook and `feasibility/Verified_Numbers.md` section 7.

Step	What we tried	Metric	Before	After	Change	Verdict
3	Personal 28-day baseline	found, r4.2 / r5.2	none	38 / 36	the bar	bar
4	Rare-site visits	found, r4.2	38	41	+3 (r5.2: 8.1)	one family
5, 6	Trained combiners	found, r4.2	39	45 to 60	+6 to +21	fail on new attacks
7	Label-free sum, job-hunters hidden	job-hunters in top 100	4 to 7	16 of 30	+9 to +12	label-free wins
8	Weekly LOF (literature)	job-hunters in top 100	5	12 to 13	+7 to +8	slow drift
10	Own-history weekly percentile	found, r4.2	38	39.4, P@20 1.00	+1.4	core
11	Family-first Stouffer fusion	found, r4.2	38	45	+7	frozen
12	Same pipeline on r5.2, frozen	found, r5.2	36	38	+2	survives
13	Snooper persistence signal	snoopers in top 100	10.9	27 of 30	+16	view
14	Exactly 100 reviews	insiders, r4.2 / r5.2	44 / 39	55 / 43	+11 / +4	default ranking
16	Recommender, tuned without labels	next-site recall@20, r4.2	0.145 (popularity)	0.211	1.45x	real job
17	Least-privilege advisor	requests per person-month	2.31 / 3.07	1.82 / 2.48	-21% / -19%	build
18	Evidence rule for Granite	hit@3, flagged insiders	0.93 / 0.89	0.98 / 0.98	+0.05 / +0.09	build
19	Early warning, day by day	flagged in time, r4.2 / r5.2	33 / 28	59 / 80	+26 / +52	headline
20	Analyst feedback; richer recommenders	insiders in 100 reviews	55 / 43	57 / 42	noise	dropped

Steps 1, 2, 9 and 15 were checks or were superseded:

- Step 1: data loading.
- Step 2: the leakage check.
- Step 9: recommender tweaks (peer-of-item: 17 found).
- Step 15: the untuned recommender verdict, replaced by Step 16.

Who builds what

Fill in names. Hours count from the start of our build.

Role	Hours 0 to 6	Hours 6 to 12	Hours 12 to 18	Hours 18 to 24
R1 Data and platform (_____)	VM: Java 17, PySpark (no pyarrow on s390x), Docker; download r4.2. At hour 2, hand a one-month sample to R2 and R4. Spark builds 3 tables: daily counts, person-computer-day, person-site-day	r5.2 if disk allows; run the whole pipeline end to end	Docker container, GitHub repo, README	Support the demo; check the repo runs from scratch
R2 Detection (_____)	Write the daily replay against the sample (`early_warning_simple.py`, 49 lines)	Run on r4.2 with the fixed settings below	r5.2 unchanged; results table vs baseline and chance, per attack type, analyst load	Export daily lists and per-person timelines for the dashboard
R3 Granite (_____)	Both IBM labs (Granite via llama.cpp; the model file must be big-endian)	Evidence pack: top 5 items per person with an ID, the person's own normal and dates; pseudonyms; prompt that must cite IDs and treats log text as data; number checker; weak-evidence flag	Pre-generate case files for the top 20; keep one live run	Case-file panel in the dashboard
R4 Recommender and demo (_____)	Dashboard skeleton: daily list, person page, case-file panel	Least-privilege advisor: implicit ALS on person-site data, pre-approve the top 20 per person, new-to-company block; count requests and blocked malicious visits	Wire everything into the dashboard	Video with on-screen proof: `uname -m` (s390x), `docker ps`, the server prompt

Fixed detector settings (decide nothing new):

- A feature is unusual when the last 7 days beat every one of that person's training weeks.
- The running score halves every 7 days.
- The list holds 20 names a day.

Recommender settings: rank 64, alpha 0.5, regularization 1.0 on log(1 + active days), using the numpy version in `recommender_als_simple.py`. Spark's ALS scales regularization differently, so if you use it, re-tune on the last training month and pick by how many next-month new sites it catches.

Cut rules

- **Hour 6:** if the Spark tables aren't ready, run r4.2 with logon, USB and file data only; add web later.
- **Hour 12:** freeze the detector settings. No tuning after this.
- **Hour 14:** if the recommender is behind, ship the new-to-company block plus each person's own history as the allow-list. Cite the recommender's local result instead.

- **Hour 18:** if Granite is slow or unreliable, show the pre-generated case files plus one live run. If it fails entirely, show the checked evidence pack itself.
- **Hour 20:** feature freeze. Only the video, README and slides from here.

Demo (3 minutes)

1. The replay plays. A leaker joins the daily list on the day of the first wikileaks upload.
2. Click the person. Granite's case file streams in, and each number turns green once the code confirms it.
3. Access panel: 20 pre-approved sites. A request for wikileaks.org is blocked as new to the company.

Say this, not that

- Say "flags insiders while they are still active", not "real-time production system". It is a replay of past logs.
- Say "the recommender predicts legitimate access", not "the recommender detects insiders". We tested that: as a detector it is no better than popularity.
- Say openly that two-day sysadmin sabotage is mostly missed, and that CERT is synthetic data.
- The early-warning method was picked after seeing results. The full detector we had named in advance also beat the baseline: 55% vs 28% of r5.2 insiders flagged in time.

Reference code (local tests; we rewrite on the server)

In `feasibility/reference/`:

- `early_warning_simple.py`: the 49-line replay behind the headline numbers.
- `early_warning_full.py`: the full replay with baseline, chance and analyst-load comparisons.
- `evidence_ranking.py`: evidence ranking for the case files.
- `recommender_als_simple.py`: the ALS recommender against popularity.
- `least_privilege_full.py`: the full least-privilege experiment.
- `org_novel_sites.py`: the new-to-company site check.

The full study is in `feasibility/OutOfRole_Feasibility.ipynb`; data download steps are in `Data_Access_Guide.pdf`.